

100
 101
 102
 103
 104
 105
 106
 107
 108
 109
 110
 111 *Image caption generation is a supervised learning*
 112 *problem generating captions based on an input image. In*
 113 *this project, I build an image captioning system that*
 114 *brings together computer vision with the natural language*
 115 *processing approach. The proposed approach consists of*
 116 *extracting visual features from images by using a pre-*
 117 *trained Convolutional Neural Network (CNN). Then the*
 118 *image features will be used to generate corresponding*
 119 *textual descriptions utilizing Recurrent Neural Networks*
 120 *(RNNs), specifically Long Short-Term Memory (LSTM)*
 121 *and Gated Recurrent Unit (GRU) models. Preprocessing*
 122 *steps are executed for both image and text data, such as*
 123 *tokenization, padding, and feature vector generation. The*
 124 *resultant system is trained to accurately predict*
 125 *informative captions for the input images, showcasing that*
 126 *integrating different deep learning methods for multi-*
 127 *modal tasks demonstrates effective results.*

Abstract

127
 128
 129
 130
 131
 132
 133
 134
 135
 136 1. Introduction
 137
 138
 139
 140
 141
 142
 143
 144
 145
 146
 147
 148
 149

Generating a caption for an image is a task located at the intersection of computer vision and natural language processing. In image captioning one generates a meaningful and grammatical sentence to describe the content of a given image. This entails recognizing visual features, as well as an understanding of color and language semantics.

A common approach to this task involves combining two models:

- A Convolutional Neural Network (CNN) to extract high-level semantic features from the image.
- A Recurrent Neural Network (RNN) for generating the caption word-by-word based on the extracted image features.

For the RNN to learn how to generate captions, it needs to be trained on annotated datasets where each image is paired with one or more descriptive sentences. During training, captions are encoded into a format that the model can understand. While one-hot encoding is a possible method, it tends to be sparse and lacks semantic richness. Instead, we use GloVe (Global Vectors for Word Representation) embeddings, which provide dense vector representations that capture semantic relationships between words more effectively.

Image Captioning with RNNs

vikrant@vt.edu

150
 151
 152
 153
 154
 155
 156
 157
 158
 159
 160
 161
 162
 163
 164
 165
 166
 167
 168
 169
 170
 171
 172
 173
 174
 175
 176
 177
 178
 179
 180
 181
 182
 183
 184
 185
 186
 187
 188
 189
 190
 191
 192
 193
 194
 195
 196
 197
 198
 199

Once training is complete, the model can generate captions by feeding it the image features (output from the CNN). The RNN then predicts the next word in the sentence iteratively until a complete caption is formed. Since images can have multiple valid descriptions, the model learns to generalize based on the diversity of captions in the training set. An example for possible captions for an image is shown in Figure 1.

To evaluate the quality of the generated captions, we use the evaluation matrix like BLEU score (Bilingual Evaluation Understudy). BLEU measures the overlap between the predicted caption and one or more ground truth references using n-gram precision.



187
 188
 189
 190
 191
 192
 193
 194
 195
 196
 197
 198
 199

Figure 1. Possible captions a.) A girl is stretched out in shallow water b.) A girl wearing a red and multi-colored bikini is laying on her back in shallow water c.) A little girl in a red swimsuit is laying on her back in shallow water d.) A young girl is lying in the sand, while ocean water is surrounding her e.) Girl wearing a bikini lying on her back in a shallow pool of clear blue water

Image captioning has wide-ranging applications, including:

- Enhancing **accessibility** for visually impaired individuals by describing images out loud.
- Improving **image search and indexing**.
- Assisting in **content creation** and summarization.
- Enabling **autonomous systems** to understand their environment in natural language.

200 201 202 203 204 205 206 207 208 209 210 211 212 213 214 215 216 217 218 219 220 221 222 223 224 225 226 227 228 229 230 231 232 233 234 235 236 237 238 239 240 241 242 243 244 245 246 247 248 249 1.1. Problem Statement

The problem statement can be summarized as building a model to generate a meaningful, grammatically correct caption with a Recurrent Neural Network (RNN)-based architecture. The model should be able to extract visual information from an image, and then guess words sequentially to generate a coherent description.

2. Approach

To solve this problem of generating meaningful caption for an Image I used an architecture, which integrated both Computer Vision using Convolution Neural Networks (CNNs) for understanding the image and Natural Language Processing using a Recurrent Neural Network (RNN) for generating caption as displayed in Figure 2.




Figure 2. Caption Generation using CNN and RNN

In the sections that follow, I describe the dataset used and the preprocessing steps applied to prepare the data for training and evaluation. I also provide a detailed explanation of the models explored and the different architectural configurations used to assess their performance.

2.1. Dataset

To train and evaluate the image captioning model, we use the Flickr8k dataset. The Flickr8k dataset consists of 8,000 images and each image have 5 unique human-annotated captions that describe the content of the image. However, for my project I used only 7,000 images that were further divided into 6,000 images for model training and remaining 1,000 images for evaluating my model.

2.2. Preprocessing

To prepare the dataset for training, we apply the following preprocessing steps:

- As the first preprocessing step, all the annotated captions are converted to lower case, removed all the punctuations and all the short or numeric word were moved.
- Then we build our vocabulary by applying a threshold for more frequent words to include only words that occurred more than 10 times. In addition to his

“startSeq” and “endSeq” were added to indicate sentence boundaries.

- Then we create our wordToIndex and IndexToWord matrix that were used for encoding word in numbers were index will start from 1 upto the maximum length of sentence in captions. These mappings are used to convert captions into numerical sequences suitable for training.
- Then we embed the words which were fed to embedding later using pre-trained GloVe embeddings (Global Vectors for Word Representation) with 200 dimensions. These embeddings provide dense vector representations of words that capture rich semantic and syntactic relationships, offering a meaningful initialization for the embedding layer in the decoder model.

2.3. Image Feature Extraction using CNN

I used a pre-trained Convolutional Neural Network (CNN) - InceptionV3 which was trained in image classification using the ImageNet dataset. The last two layers that were used for prediction were removed so that the image features could be extracted from the given image. This returns a 2048-dimensional vector representing the image's content as displayed in Figure 3.




Figure 3. Feature Vector Extraction (Feature Engineering) from InceptionV3


2.4. Caption Generation using RNN

The word embeddings generated during the preprocessing step are fed into the Recurrent Neural Network (RNN) component of the model. This is responsible for processing the sequence of embedded words and learning the temporal dependencies between them — essentially modeling how words are structured in a grammatically and semantically correct sentence. To explore the effectiveness of different sequence modeling techniques, I experimented with two widely used RNN architectures:

- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

The different configurations of the image captioning model explored in this project are described below:

300
301
302
303 Configurations 1: This is the baseline model configuration
304 where a single-layer LSTM is used for predictions with
305 256 neurons. The model was trained using the Adam
306 optimizer, which is effective for achieving faster and more
307 stable convergence. The model was trained for 20 epochs,
308 as shown in *Figure 4*.




329 *Figure 4. Baseline Model using LSTM*

330 The choice of 20 epochs was based on observation — both
331 the training loss and training accuracy stabilized by this
332 point, indicating that the model had learned meaningful
333 patterns without overfitting. Since the task is to predict the
334 next word from a vocabulary of possible words at each
335 time step, this is treated as a multi-class classification
336 problem. Therefore, categorical cross-entropy was used as
337 the loss function.

338 Configurations 2: To explore the performance of an
339 alternative RNN architecture, I implemented the same
340 encoder-decoder setup using a Gated Recurrent Unit
341 (GRU) with 256 hidden units as displayed in *Figure 5*.
342 GRUs are known for their computational efficiency and
343 ability to capture temporal dependencies in sequences,
344 often providing performance comparable to LSTMs with
345 fewer parameters.

346 Unlike the previous configuration that used the Adam
347 optimizer, this model was trained using Stochastic
348 Gradient Descent (SGD) as it updates the weights on small
349 batches of data. While it may converge more slowly than
adaptive optimizers like Adam, SGD offers greater control
through its learning rate and often leads to better
generalization. The model was trained for 20 epochs and
the same architecture and preprocessing steps as

350 Configuration 1 were retained to enable a fair comparison.
351 The categorical cross-entropy loss function was again
352 used, as the task remains a multi-class prediction problem.
353



355 *Figure 5. Baseline Model using GRU*

356 Configurations 3: The baseline GRU model illustrated a
357 considerably poorer performance than the baseline LSTM
358 model, as anticipated because GRUs are lighter and faster,
359 but often have difficulty maintaining longer-range
360 dependencies in sequences compared to LSTMs. To help
361 mitigate this limitation, I created a stacked GRU
362 architecture with three GRU layers, each with 256 units,
363 which will increase the model's ability to capture complex
364 temporal dependencies and hold more context over longer
365 sequences. The architecture of this model is shown in
366 Figure 6.

367 To mitigate the risk of overfitting and improve
368 generalization, dropout with a rate of 0.5 was used
369 between the GRU layers. Dropout helps with
370 regularization of the network by randomly dropping units
371 during training so that the model learns more robust
372 features. This model was also trained with SGD initially,
373 as in Configuration 2, but I immediately recognized that,
374 due to the increased deepness and complexity, the model
375 was not converging well with SGD. I then switched to the
376 RMSprop optimizer, which adapts the learning rate for
377 each parameter and is useful for training deeper recurrent
378 networks. The model was trained for 20 epochs, as with
379 the other configurations, and categorical cross entropy was
380 selected as the loss function, appropriate for the multiclass
381 prediction.




Figure 6. Stacked Model with 3 GRU Layers

Configurations 4: After realizing the significant improvement in performance of my stacked GRU model I extended the same concept to the LSTM architecture by implementing a stacked LSTM model with three LSTM layers, each containing 256 units. The motivation was to leverage LSTM's superior ability to retain long-term dependencies and evaluate whether deeper architectures could further enhance caption quality. The architecture of this model is illustrated in Figure 7. To mitigate the risk of overfitting associated with deep LSTM networks, I added Dropout layers with a rate of 0.5 between each LSTM layer.

For optimization, I continued using the Adam optimizer, known for its efficiency and adaptive learning capabilities. However, to exert finer control over the training process and reduce the chance of overfitting in this deeper model, I manually adjusted the learning rate to 0.001. This adjustment provided more stable and gradual convergence. Despite these modifications, the performance improvement over the baseline LSTM was relatively modest. This suggests that while LSTM networks

inherently handle long-term dependencies well, simply stacking more layers does not guarantee significantly better results — especially on smaller datasets like Flickr8k, where deep models may not generalize effectively. This configuration was essential in understanding the trade-offs between model depth and overfitting, and it provided a deeper insight into how LSTM behavior scales with complexity.



Figure 7. Stacked Model with 3 LSTM Layers

2.5. Output layers

The final stage of the model involves adding the image features extracted by the CNN with the contextual information generated by the RNN using an add() layer. The merged vector is then passed through a fully connected (Dense) output layer, which produces a probability distribution over the vocabulary. Based on this model predicts the most likely next word in the sequence, given the image context and the previously generated words.

500	550
501	551
502	552
503	553
504	554
505	555
506	556
507	557
508	558
509	559
510	560
511	561
512	562
513	563
514	564
515	565
516	566
517	567
518	568
519	569
520	570
521	571
522	572
523	573
524	574
525	575
526	576
527	577
528	578
529	579
530	580
531	581
532	582
533	583
534	584
535	585
536	586
537	587
538	588
539	589
540	590
541	591
542	592
543	593
544	594
545	595
546	596
547	597
548	598
549	599

2.6. Training

To train the model, I provided the model with the image features as well as the caption sequences. Since the feature vector (image + caption input) can be large, especially with high-dimensional image embeddings and word sequences, it is not feasible to load the entire dataset into memory at the same time.

To address this issue, the training data is divided into mini batches. Each batch consists of:

- The image feature vector extracted by the CNN (InceptionV3)
- The partial caption sequence up to the current word
- And the target word the model should predict next.

2.7. Evaluation

I employed the BLEU score to assess the quality of the generated captions. The BLEU score is a conventional comparison of the generated caption to one or more reference captions, based on n-gram precision (1-gram to 4-grams), for all cases of the models.

3. Experimental Results

After training each model configuration, I evaluated their performance by generating captions for unseen images from the test set.

3.1. Predicted Caption with Images

The following examples from each model show the predicted caption alongside the input image that was presented to the trained model. These examples are useful in determining how well each architecture is modeling the visual input and extracting meaning in language. All models were run on the same set of images for comparison here. All predicted captions indicate the model's understanding of objects, actions, and relationships in the image.




Figure 8. Generated Caption from base LSTM model



Figure 9. Generated Caption from base LSTM model using Greedy Search



Figure 10. Generated Caption from base GRU model



Figure 11. Generated Caption from stacked GRU model



Figure 12. Generated Caption from stacked LSTM model




Figure 14. Training Loss and Accuracy for Base GRU model




Figure 15. Training Loss and Accuracy for Stacked LSTM model




Figure 16. Training Loss and Accuracy for Stacked GRU model

4. Discussion

4.1 Predicted Captions Results

When generating the caption, I considered two decoding approaches:

- Greedy decoding, where the highest probability word is selected at each time step.
- A manual, step-by-step, prediction loop, allowing for more control of the word generation and debugging.

For the Baseline LSTM model, I first tried greedy decoding. I noticed that the model was generating a lot of repeated outputs, such as predicting the same word multiple times in a row as shown in Table 1 and Figure 9. This was a sign that the model had not generalized sufficiently and was too confident when predicting certain tokens. The repeated predictions were likely due to the model being limited in learning a range of sequence transitions. Because of these issues, I decided not to use greedy decoding in any of the models. Instead, I had a custom prediction method, where I could generate

captions word by word and could observe the model outputs in more detail. This allowed for more stable and interpretable results across the configurations. After this I generated captions for each of the four model configurations, I was able to observe some notable differences in the output quality, structure and overall accuracy for context that discussed below:

- Model 1: For Model 1, which used a single-layer LSTM with 256 units, the model was able to produce captions that were generally grammatically correct and coherent. The captions that it predicted were constantly between 5 to 8 words, which aligns well with the average caption length across the training data, meaning it was able to learn the rhythmic nature of sentence structures in the training data. However, the model occasionally had specific semantic opportunities for improvement. For example, in images that featured a single dog without a collar, the model tended to predict, "a dog with a black collar". These errors occur due to training bias or overfitting — the model had learned to correlate frequently identified patterns (dog, collar) regardless of whether the collar was present in that instance. In conclusion, while the baseline LSTM demonstrated a good ability to model sentence structures and predict frequent visual instances (dog with collar), it did sometimes hallucinate details that the image had not visually specified.

- Model 2: The results of Model 2 indicated that having a single-layer Gated Recurrent Unit (GRU) with 256 neurons didn't perform well compared to the LSTM baseline. In many cases, the model failed to generate suitable caption-words at all: it either returned with empty captions or repeated the same word, e.g., "dog dog dog dog...". Repetition and lack-of-sentences is reflective of the GRU's short capacity for long-term context retention when used in shallow setup. This can also indicate that the model struggled to learn the temporal aspects of generating words in order to learn how to form sentences during training. Overall, the baseline GRU lacked the stability and language signature understanding of the LSTM and could not generalize across different image contexts.

- Model 3: Model 3 represented an important upgrade from the baseline GRU structure. For this model, I stacked three GRU layers of 256 units each and added a dropout rate of 0.5 to reduce overfitting. Another important change was switching the SGD optimizer for RMSprop, which is ideally suited to training deeper recurrent networks based on its adaptive learning. Overall, this architecture led to an observable improvement in performance. In contrast

800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849

to the GRU baseline, the stacked GRU was able to produce longer, increasingly cohesive, and contextually accurate captions, comparable to the baseline LSTM. That said, while there was generally improved performance, the model still displayed certain obvious inconsistencies. For example, the model had a tendency to begin with "two" regardless of whether or not multiple subjects were depicted in the image, hinting that the model had learned structural conventions well, but was still prone semantic hallucinations or the biases of the dataset. Overall, the fairly strong performance of this stacked GRU architecture warranted continuing exploration of deeper LSTM architectures with the potential to achieve even greater improvement improved accuracy and stability of captioning results.

- Model 4: This model set up was the overall best-performing in terms of fluency and semantic coherence. For Model 4, I stacked three LSTM layers of 256 units and used a dropout of 0.5 between layers for regularization and to lessen overfitting. The model was trained using the Adam optimizer at a learning rate of 0.001 for stable and efficient convergence. The generated predictions from the model were generally correct, fluent, and semantically correct, showing improvement in contextual understanding and diminished repetition compared to prior models. Though improvements were made in the LSTM architecture, the gain in improvement was minimal compared to the baseline LSTM. These observations indicate that more complex and deeper LSTM architectures can improve the model capacity, but the scope in performance gain experienced was not as substantial as that seen moving from the baseline GRU to a stacked GRU, which potentially also reflects the restrictions in dataset size (Flickr8k) in which more complexity in the model does not lead to greater generalization.

4.2 Predicted Captions Results

BLEU (Bilingual Evaluation Understudy) score is a popular evaluation metric for natural language generation tasks, including machine translation and image captioning. It quantifies the similarity between a sentence produced by a model (the hypothesis) and one or more reference sentences produced by humans. I used this metric to evaluate the model and below were the results:

- The **Stacked LSTM model** achieved the highest scores across all BLEU metrics, indicating the strongest overall caption generation quality.
- The Base GRU model showed very weak performance, with a BLEU-1 score nearing zero and a zero score for BLEU-2, BLEU-3, and BLEU-4. This indicates that the model hardly produced meaningful n-gram

850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899

sequences and was unable to match up to even simple bigram structures from the reference captions.

- The low scores can be attributed to both the limited learning capacity of a shallow GRU and the use of the SGD optimizer, which struggled to converge effectively.
- The Stacked GRU model, by contrast, showed major improvements, particularly in higher-order BLEU scores, suggesting that depth and better optimization (RMSprop) helped it retain longer-term context.
- The Base LSTM provided a solid baseline, while the Stacked LSTM slightly improved upon it — indicating that LSTM models remain more stable and expressive across both shallow and deep configurations.

4.3 Training Loss and Accuracy

- Model 1: For the model 1 with base LSTM, the continual uptrend in accuracy and downtrend in loss indicates that the LSTM base model was both well-optimized and stable. The curves both flatten in the end, indicating that training for more than 20 epochs would produce minimal gains unless adjustments were made to learning rates or monitoring validation.
- Model 2: The GRU base model did not reach any significant learning during training, reflecting the very low to zero BLEU scores we obtained. The poor performance is likely a result of:
 - The shallow GRU architecture (single layer and shallow capacity).
 - The SGD optimizer

The accuracy curve and loss both indicate the model likely reached a learning ceiling quite quickly and was not able to generalize well.

- Model 3: The stacked GRU model exhibited considerable advancements compared to the GRU baseline, which presented almost zero BLEU scores and relatively constant training curves. While the same forms of training were observed, there were no signs of overfitting, and the training process was uniform and smooth enough that we can conclude that RMSprop was indeed a significantly better choice of optimizer for GRU over SGD. While the stacked GRU approach did not perform as well as the LSTM models, it is clear that the stacked GRU closes the performance gap enough that it is still a reasonable alternative when configured appropriately.

- Model 4: The training curves for the stacked LSTM are very similar to the base LSTM, but they appear to have slightly more stable and slower you increase, because of the increased depth of the model and lower learning rate. Despite the size of the model, there are no signs of overfitting in the curves, likely due to the use of dropout layers and tuning of the learning rate.

900	950
901	951
902	952
903	953
904	954
905	955
906	956
907	957
908	958
909	959
910	960
911	961
912	962
913	963
914	964
915	965
916	966
917	967
918	968
919	969
920	970
921	971
922	972
923	973
924	974
925	975
926	976
927	977
928	978
929	979
930	980
931	981
932	982
933	983
934	984
935	985
936	986
937	987
938	988
939	989
940	990
941	991
942	992
943	993
944	994
945	995
946	996
947	997
948	998
949	999
Overall, the minor improvement in accuracy and in loss compared to the base LSTM suggests that stacking LSTMs doesn't inherently improve performance - a finding supported by the BLEU scores.	
5. Conclusion	
In conclusion during this work I investigated the task of generating image captions using a standard encoder-decoder framework with CNNs used for image feature extraction and RNNs used for generating sequences of captions. I learned how different choices of model depth, optimizer, and regularization impact the quality of the generated captions, a series of experiments were conducted with LSTM and GRU models, in both baseline and stacked configurations. As a result of the experiments, it is clear that the model depth, optimizer choice, and regularization are all important decision points influencing the quality of the generated captions in some way. Results indicate that the best performance overall came from the stacked LSTM model, while the stacked GRU model produced competitive results (with correct configuration). The research shows how model architectures and strategic training can greatly affect performance when building multimodal language models.	
References	
[1] https://medium.com/data-science/image-captioning-with-keras-teaching-computers-to-describe-pictures-c88a46a311b8	
[2] https://www.digitalocean.com/community/tutorials/bleu-score-in-python	
[3] https://www.kaggle.com/datasets/adityajn105/flickr8k	
[4] https://learn.microsoft.com/en-us/azure/ai-services/translator/custom-translator/concepts/bleu-score	